

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

import torch
import torch.optim as optim
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import Dataset, DataLoader

import cv2
import albumentations as A
from albumentations.pytorch import ToTensorV2

import copy
import os
import random

```

Подготовка данных

```

TRAIN_PATH = '/kaggle/input/tinkofftestccpd2019dl1/CCPD2019-dl1/train'
TEST_PATH = '/kaggle/input/tinkofftestccpd2019dl1/CCPD2019-dl1/test'

```

```

def get_label_chars_set():
    all_chars = set()
    for filename in os.listdir(TRAIN_PATH):
        label = filename.split('-')[-1][:-4]
        for char in label:
            all_chars.add(char)

    for filename in os.listdir(TEST_PATH):
        label = filename.split('-')[-1][:-4]
        for char in label:
            all_chars.add(char)
    return list(all_chars)

class CarNumsDataset(Dataset):
    def __init__(self, image_dir, image_names, transform=None):
        self.image_dir = image_dir
        self.image_names = image_names
        self.transform = transform
        self.label_chars_dict = {char : i for i, char in
enumerate(get_label_chars_set())}

    def __len__(self):
        return len(self.image_names)

    def __getitem__(self, idx):
        image_path = os.path.join(self.image_dir,
self.image_names[idx])

```

```

        image = cv2.imread(image_path)
        image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
        label = self.image_names[idx].split('-')[-1][:4]
        label = self._transform_label(label)

        if self.transform is not None:
            image = self.transform(image=image)['image']

        return image, label

    def _transform_label(self, label):
        label_vector = [self.label_chars_dict[char] for char in label]
        return torch.tensor(label_vector, dtype=torch.long)

DATASET_MEAN = [0.2656, 0.4230, 0.6284]
DATASET_STD = [0.2096, 0.2029, 0.2181]

train_transform = A.Compose(
    [
        A.PixelDropout(dropout_prob=0.05),
        A.ShiftScaleRotate(shift_limit=0.0, scale_limit=0.01,
rotate_limit=5, p=0.5),
        A.RGBShift(r_shift_limit=15, g_shift_limit=15,
b_shift_limit=15, p=0.5),
        A.RandomBrightnessContrast(p=0.5),
        A.augmentations.geometric.Resize(128, 256),
        A.Normalize(mean=DATASET_MEAN, std=DATASET_STD),
        ToTensorV2()
    ]
)

val_transform = A.Compose(
    [
        A.augmentations.geometric.Resize(128, 256),
        A.Normalize(mean=DATASET_MEAN, std=DATASET_STD),
        ToTensorV2()
    ]
)

test_transform = A.Compose(
    [
        A.augmentations.geometric.Resize(128, 256),
        A.Normalize(mean=DATASET_MEAN, std=DATASET_STD),
        ToTensorV2()
    ]
)

def visualize_augmentations(dataset, idx=0, samples=10, cols=5):
    dataset = copy.deepcopy(dataset)
    dataset.transform = A.Compose([t for t in dataset.transform if not

```

```

isinstance(t, (A.Normalize, ToTensorV2)))]
    rows = samples // cols
    figure, ax = plt.subplots(nrows=rows, ncols=cols, figsize=(12, 6))
    for i in range(samples):
        image, _ = dataset[idx]
        ax.ravel()[i].imshow(image)
        ax.ravel()[i].set_axis_off()
    plt.tight_layout()
    plt.show()

train_images = os.listdir(TRAIN_PATH)
train_images.sort()
random.seed(0)
random.shuffle(train_images) # for random train/val split further
test_images = os.listdir(TEST_PATH)
test_images.sort()

train_size = int(0.8 * len(train_images))
train_image_names = train_images[:train_size]
val_image_names = train_images[train_size:]
test_image_names = test_images

train_dataset = CarNumsDataset(TRAIN_PATH, train_image_names,
train_transform)
val_dataset = CarNumsDataset(TRAIN_PATH, val_image_names,
val_transform)
test_dataset = CarNumsDataset(TEST_PATH, test_image_names,
test_transform)

train_dataloader = DataLoader(train_dataset, batch_size=256,
num_workers=2, shuffle=True)
val_dataloader = DataLoader(val_dataset, batch_size=256,
num_workers=2)
test_dataloader = DataLoader(test_dataset, batch_size=256)

visualize_augmentations(train_dataset, 25)

```



Создание и обучение модели

```
DEVICE = torch.device('cuda') if torch.cuda.is_available() else
torch.device('cpu')
DEVICE

device(type='cuda')

from torchvision.models import resnet18

class OCRModel(nn.Module):
    def __init__(self, n_classes):
        super().__init__()
        self.resnet =
nn.Sequential(*list(resnet18(pretrained=False).children())[:-2])
        self.conv1 = nn.Conv1d(in_channels=512,
out_channels=n_classes, kernel_size=2,
padding=0)

    def forward(self, x):
        x = self.resnet(x)
        x = x.mean(dim=2)
        x = self.conv1(x)
        return x

def fit(epochs, model, optimizer, loss_fn):
    for epoch in range(1, epochs + 1):
        print(f'EPOCH_{epoch} START')

        train_loss = 0
        val_loss = 0
        train_acc = 0
        val_acc = 0
        train_cer = 0
        val_cer = 0

        model.train()
        for i, (xb, yb) in enumerate(train_dataloader):
            xb, yb = xb.to(DEVICE), yb.to(DEVICE)
            yb_logits = model(xb)

            loss = loss_fn(yb_logits, yb)
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

            train_loss += loss.detach().cpu().item()
            yb_preds = yb_logits.argmax(dim=1) #[bs, 66 , 7] -> [bs,
7]

            acc = ((yb_preds == yb).all(dim=1)).float().mean()
            cer = (yb_preds != yb).float().mean()
```

```

train_acc += acc.detach().cpu().item()
train_cer += cer.detach().cpu().item()

if i % 100 == 0:
    batch_num = i + 1
    avg_loss = train_loss / batch_num
    avg_acc = train_acc / batch_num
    avg_cer = train_cer / batch_num
    print(f'TRAIN_BATCH_NUM: {batch_num} TRAIN_LOSS:
{avg_loss:.5f} TRAIN_ACC: {avg_acc:.5f} TRAIN_CER: {avg_cer:.5f}')

train_loss /= len(train_dataloader)
train_acc /= len(train_dataloader)
train_cer /= len(train_dataloader)

model.eval()
with torch.inference_mode():
    for xb, yb in val_dataloader:
        xb, yb = xb.to(DEVICE), yb.to(DEVICE)
        yb_logits = model(xb)

        loss = loss_fn(yb_logits, yb)
        val_loss += loss.detach().cpu().item()
        yb_preds = yb_logits.argmax(dim=1)
        acc = ((yb_preds == yb).all(dim=1)).float().mean()
        cer = (yb_preds != yb).float().mean()
        val_acc += acc.detach().cpu().item()
        val_cer += cer.detach().cpu().item()

val_loss /= len(val_dataloader)
val_acc /= len(val_dataloader)
val_cer /= len(val_dataloader)

print(f'EPOCH_{epoch} END')
print(f'TRAIN_LOSS: {train_loss:.5f} TRAIN_ACC:
{train_acc:.5f} TRAIN_CER: {train_cer:.5f}')
print(f'VAL_LOSS: {val_loss:.5f} VAL_ACC: {val_acc:.5f}
VAL_CER: {val_cer:.5f}')
print('-'*50)

n_classes = len(get_label_chars_set())
model = OCRModel(n_classes=n_classes)
model.to(DEVICE)
loss_fn = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)
epochs = 10

fit(epochs, model, optimizer, loss_fn)

```

```
EPOCH_1 START
TRAIN_BATCH_NUM: 1 TRAIN_LOSS: 4.35284 TRAIN_ACC: 0.00000 TRAIN_CER:
0.99442
TRAIN_BATCH_NUM: 101 TRAIN_LOSS: 0.35784 TRAIN_ACC: 0.73886 TRAIN_CER:
0.08341
TRAIN_BATCH_NUM: 201 TRAIN_LOSS: 0.19191 TRAIN_ACC: 0.85236 TRAIN_CER:
0.04463
TRAIN_BATCH_NUM: 301 TRAIN_LOSS: 0.13290 TRAIN_ACC: 0.89458 TRAIN_CER:
0.03089
TRAIN_BATCH_NUM: 401 TRAIN_LOSS: 0.10294 TRAIN_ACC: 0.91628 TRAIN_CER:
0.02394
TRAIN_BATCH_NUM: 501 TRAIN_LOSS: 0.08390 TRAIN_ACC: 0.93062 TRAIN_CER:
0.01953
TRAIN_BATCH_NUM: 601 TRAIN_LOSS: 0.07137 TRAIN_ACC: 0.93981 TRAIN_CER:
0.01664
EPOCH_1 END
TRAIN_LOSS: 0.06890 TRAIN_ACC: 0.94176 TRAIN_CER: 0.01606
VAL_LOSS: 0.00444 VAL_ACC: 0.99341 VAL_CER: 0.00101
-----
EPOCH_2 START
TRAIN_BATCH_NUM: 1 TRAIN_LOSS: 0.00630 TRAIN_ACC: 0.99219 TRAIN_CER:
0.00112
TRAIN_BATCH_NUM: 101 TRAIN_LOSS: 0.00632 TRAIN_ACC: 0.99022 TRAIN_CER:
0.00161
TRAIN_BATCH_NUM: 201 TRAIN_LOSS: 0.00565 TRAIN_ACC: 0.99083 TRAIN_CER:
0.00149
TRAIN_BATCH_NUM: 301 TRAIN_LOSS: 0.00552 TRAIN_ACC: 0.99089 TRAIN_CER:
0.00145
TRAIN_BATCH_NUM: 401 TRAIN_LOSS: 0.00560 TRAIN_ACC: 0.99114 TRAIN_CER:
0.00143
TRAIN_BATCH_NUM: 501 TRAIN_LOSS: 0.00547 TRAIN_ACC: 0.99124 TRAIN_CER:
0.00141
TRAIN_BATCH_NUM: 601 TRAIN_LOSS: 0.00532 TRAIN_ACC: 0.99156 TRAIN_CER:
0.00136
EPOCH_2 END
TRAIN_LOSS: 0.00533 TRAIN_ACC: 0.99156 TRAIN_CER: 0.00136
VAL_LOSS: 0.00371 VAL_ACC: 0.99398 VAL_CER: 0.00092
-----
EPOCH_3 START
TRAIN_BATCH_NUM: 1 TRAIN_LOSS: 0.00081 TRAIN_ACC: 1.00000 TRAIN_CER:
0.00000
TRAIN_BATCH_NUM: 101 TRAIN_LOSS: 0.00398 TRAIN_ACC: 0.99350 TRAIN_CER:
0.00105
TRAIN_BATCH_NUM: 201 TRAIN_LOSS: 0.00350 TRAIN_ACC: 0.99413 TRAIN_CER:
0.00095
TRAIN_BATCH_NUM: 301 TRAIN_LOSS: 0.00353 TRAIN_ACC: 0.99398 TRAIN_CER:
0.00095
TRAIN_BATCH_NUM: 401 TRAIN_LOSS: 0.00348 TRAIN_ACC: 0.99402 TRAIN_CER:
0.00094
TRAIN_BATCH_NUM: 501 TRAIN_LOSS: 0.00360 TRAIN_ACC: 0.99374 TRAIN_CER:
```

```
0.00098
TRAIN_BATCH_NUM: 601 TRAIN_LOSS: 0.00353 TRAIN_ACC: 0.99388 TRAIN_CER:
0.00095
EPOCH_3 END
TRAIN_LOSS: 0.00351 TRAIN_ACC: 0.99392 TRAIN_CER: 0.00095
VAL_LOSS: 0.00245 VAL_ACC: 0.99575 VAL_CER: 0.00064
-----
EPOCH_4 START
TRAIN_BATCH_NUM: 1 TRAIN_LOSS: 0.00127 TRAIN_ACC: 1.00000 TRAIN_CER:
0.00000
TRAIN_BATCH_NUM: 101 TRAIN_LOSS: 0.00209 TRAIN_ACC: 0.99613 TRAIN_CER:
0.00057
TRAIN_BATCH_NUM: 201 TRAIN_LOSS: 0.00225 TRAIN_ACC: 0.99590 TRAIN_CER:
0.00061
TRAIN_BATCH_NUM: 301 TRAIN_LOSS: 0.00272 TRAIN_ACC: 0.99499 TRAIN_CER:
0.00076
TRAIN_BATCH_NUM: 401 TRAIN_LOSS: 0.00270 TRAIN_ACC: 0.99517 TRAIN_CER:
0.00074
TRAIN_BATCH_NUM: 501 TRAIN_LOSS: 0.00265 TRAIN_ACC: 0.99529 TRAIN_CER:
0.00072
TRAIN_BATCH_NUM: 601 TRAIN_LOSS: 0.00262 TRAIN_ACC: 0.99530 TRAIN_CER:
0.00072
EPOCH_4 END
TRAIN_LOSS: 0.00267 TRAIN_ACC: 0.99524 TRAIN_CER: 0.00073
VAL_LOSS: 0.00349 VAL_ACC: 0.99458 VAL_CER: 0.00082
-----
EPOCH_5 START
TRAIN_BATCH_NUM: 1 TRAIN_LOSS: 0.00532 TRAIN_ACC: 0.99219 TRAIN_CER:
0.00112
TRAIN_BATCH_NUM: 101 TRAIN_LOSS: 0.00270 TRAIN_ACC: 0.99505 TRAIN_CER:
0.00080
TRAIN_BATCH_NUM: 201 TRAIN_LOSS: 0.00254 TRAIN_ACC: 0.99559 TRAIN_CER:
0.00071
TRAIN_BATCH_NUM: 301 TRAIN_LOSS: 0.00265 TRAIN_ACC: 0.99547 TRAIN_CER:
0.00072
TRAIN_BATCH_NUM: 401 TRAIN_LOSS: 0.00277 TRAIN_ACC: 0.99519 TRAIN_CER:
0.00076
TRAIN_BATCH_NUM: 501 TRAIN_LOSS: 0.00264 TRAIN_ACC: 0.99537 TRAIN_CER:
0.00074
TRAIN_BATCH_NUM: 601 TRAIN_LOSS: 0.00268 TRAIN_ACC: 0.99531 TRAIN_CER:
0.00074
EPOCH_5 END
TRAIN_LOSS: 0.00267 TRAIN_ACC: 0.99533 TRAIN_CER: 0.00074
VAL_LOSS: 0.00215 VAL_ACC: 0.99684 VAL_CER: 0.00050
-----
EPOCH_6 START
TRAIN_BATCH_NUM: 1 TRAIN_LOSS: 0.00010 TRAIN_ACC: 1.00000 TRAIN_CER:
0.00000
TRAIN_BATCH_NUM: 101 TRAIN_LOSS: 0.00201 TRAIN_ACC: 0.99675 TRAIN_CER:
0.00049
```

TRAIN_BATCH_NUM: 201 TRAIN_LOSS: 0.00180 TRAIN_ACC: 0.99675 TRAIN_CER:
0.00048
TRAIN_BATCH_NUM: 301 TRAIN_LOSS: 0.00179 TRAIN_ACC: 0.99689 TRAIN_CER:
0.00047
TRAIN_BATCH_NUM: 401 TRAIN_LOSS: 0.00194 TRAIN_ACC: 0.99650 TRAIN_CER:
0.00053
TRAIN_BATCH_NUM: 501 TRAIN_LOSS: 0.00189 TRAIN_ACC: 0.99665 TRAIN_CER:
0.00051
TRAIN_BATCH_NUM: 601 TRAIN_LOSS: 0.00196 TRAIN_ACC: 0.99650 TRAIN_CER:
0.00054

EPOCH_6 END

TRAIN_LOSS: 0.00196 TRAIN_ACC: 0.99647 TRAIN_CER: 0.00054
VAL_LOSS: 0.00221 VAL_ACC: 0.99669 VAL_CER: 0.00051

EPOCH_7 START

TRAIN_BATCH_NUM: 1 TRAIN_LOSS: 0.00045 TRAIN_ACC: 1.00000 TRAIN_CER:
0.00000
TRAIN_BATCH_NUM: 101 TRAIN_LOSS: 0.00153 TRAIN_ACC: 0.99714 TRAIN_CER:
0.00041
TRAIN_BATCH_NUM: 201 TRAIN_LOSS: 0.00201 TRAIN_ACC: 0.99652 TRAIN_CER:
0.00053
TRAIN_BATCH_NUM: 301 TRAIN_LOSS: 0.00191 TRAIN_ACC: 0.99678 TRAIN_CER:
0.00049
TRAIN_BATCH_NUM: 401 TRAIN_LOSS: 0.00182 TRAIN_ACC: 0.99684 TRAIN_CER:
0.00048
TRAIN_BATCH_NUM: 501 TRAIN_LOSS: 0.00178 TRAIN_ACC: 0.99696 TRAIN_CER:
0.00047
TRAIN_BATCH_NUM: 601 TRAIN_LOSS: 0.00175 TRAIN_ACC: 0.99691 TRAIN_CER:
0.00047

EPOCH_7 END

TRAIN_LOSS: 0.00179 TRAIN_ACC: 0.99687 TRAIN_CER: 0.00048
VAL_LOSS: 0.00217 VAL_ACC: 0.99647 VAL_CER: 0.00053

EPOCH_8 START

TRAIN_BATCH_NUM: 1 TRAIN_LOSS: 0.00168 TRAIN_ACC: 0.99609 TRAIN_CER:
0.00056
TRAIN_BATCH_NUM: 101 TRAIN_LOSS: 0.00169 TRAIN_ACC: 0.99675 TRAIN_CER:
0.00047
TRAIN_BATCH_NUM: 201 TRAIN_LOSS: 0.00190 TRAIN_ACC: 0.99619 TRAIN_CER:
0.00056
TRAIN_BATCH_NUM: 301 TRAIN_LOSS: 0.00186 TRAIN_ACC: 0.99630 TRAIN_CER:
0.00055
TRAIN_BATCH_NUM: 401 TRAIN_LOSS: 0.00176 TRAIN_ACC: 0.99662 TRAIN_CER:
0.00051
TRAIN_BATCH_NUM: 501 TRAIN_LOSS: 0.00188 TRAIN_ACC: 0.99639 TRAIN_CER:
0.00054
TRAIN_BATCH_NUM: 601 TRAIN_LOSS: 0.00198 TRAIN_ACC: 0.99628 TRAIN_CER:
0.00056

EPOCH_8 END

TRAIN_LOSS: 0.00202 TRAIN_ACC: 0.99624 TRAIN_CER: 0.00057


```

VAL_LOSS:    0.00390 VAL_ACC:    0.99358 VAL_CER:    0.00097
-----
EPOCH_9 START
TRAIN_BATCH_NUM: 1 TRAIN_LOSS: 0.00313 TRAIN_ACC: 1.00000 TRAIN_CER:
0.00000
TRAIN_BATCH_NUM: 101 TRAIN_LOSS: 0.00115 TRAIN_ACC: 0.99768 TRAIN_CER:
0.00035
TRAIN_BATCH_NUM: 201 TRAIN_LOSS: 0.00168 TRAIN_ACC: 0.99689 TRAIN_CER:
0.00048
TRAIN_BATCH_NUM: 301 TRAIN_LOSS: 0.00161 TRAIN_ACC: 0.99704 TRAIN_CER:
0.00046
TRAIN_BATCH_NUM: 401 TRAIN_LOSS: 0.00157 TRAIN_ACC: 0.99713 TRAIN_CER:
0.00046
TRAIN_BATCH_NUM: 501 TRAIN_LOSS: 0.00147 TRAIN_ACC: 0.99726 TRAIN_CER:
0.00043
TRAIN_BATCH_NUM: 601 TRAIN_LOSS: 0.00156 TRAIN_ACC: 0.99717 TRAIN_CER:
0.00045
EPOCH_9 END
TRAIN_LOSS: 0.00156 TRAIN_ACC: 0.99714 TRAIN_CER: 0.00045
VAL_LOSS:    0.00182 VAL_ACC:    0.99756 VAL_CER:    0.00038
-----
EPOCH_10 START
TRAIN_BATCH_NUM: 1 TRAIN_LOSS: 0.00003 TRAIN_ACC: 1.00000 TRAIN_CER:
0.00000
TRAIN_BATCH_NUM: 101 TRAIN_LOSS: 0.00195 TRAIN_ACC: 0.99648 TRAIN_CER:
0.00056
TRAIN_BATCH_NUM: 201 TRAIN_LOSS: 0.00158 TRAIN_ACC: 0.99708 TRAIN_CER:
0.00046
TRAIN_BATCH_NUM: 301 TRAIN_LOSS: 0.00153 TRAIN_ACC: 0.99727 TRAIN_CER:
0.00043
TRAIN_BATCH_NUM: 401 TRAIN_LOSS: 0.00152 TRAIN_ACC: 0.99733 TRAIN_CER:
0.00042
TRAIN_BATCH_NUM: 501 TRAIN_LOSS: 0.00152 TRAIN_ACC: 0.99736 TRAIN_CER:
0.00042
TRAIN_BATCH_NUM: 601 TRAIN_LOSS: 0.00161 TRAIN_ACC: 0.99723 TRAIN_CER:
0.00044
EPOCH_10 END
TRAIN_LOSS: 0.00161 TRAIN_ACC: 0.99719 TRAIN_CER: 0.00044
VAL_LOSS:    0.00219 VAL_ACC:    0.99706 VAL_CER:    0.00044
-----

torch.save(model.state_dict(), 'model_10_epochs.pt')

```

Подсчет метрик

```

def get_model_predictions(model, dataloader):
    preds = []
    model.eval()
    for xb, yb in dataloader:
        xb, yb = xb.to(DEVICE), yb.to(DEVICE)

```

```

        yb_logits = model(xb)
        yb_preds = yb_logits.argmax(dim=1)
        preds.append(yb_preds)
    return torch.concat(preds)

def get_all_target(dataset):
    dataloader = DataLoader(dataset, batch_size=len(dataset))
    xb, yb = next(iter(dataloader))
    yb = yb.to(DEVICE)
    return yb

test_target = get_all_target(test_dataset)
test_preds = get_model_predictions(model, test_dataloader)

```

МЕТРИКИ

```

test_acc = (test_target ==
test_preds).all(dim=1).float().mean().item()
test_cer = (test_target != test_preds).float().mean().item()
print(f'TEST_ACC: {test_acc:.5f} TEST_CER: {test_cer:.4f}')

```

TEST_ACC: 0.98280 TEST_CER: 0.0028

Анализ ошибок

```

wrong_preds_idx = (~(test_target == test_preds).all(dim=1))
wrong_preds_idx = wrong_preds_idx.nonzero().squeeze().cpu().numpy()
wrong_preds_idx

```

```

array([  8,  44,  71,  94, 140, 161, 239, 438, 465, 483,
507,
      575, 627, 688, 714, 735, 737, 775, 890, 1021, 1028,
1064,
     1182, 1251, 1380, 1489, 1655, 1694, 1860, 1882, 2004, 2028,
2087,
     2177, 2274, 2279, 2351, 2612, 2775, 2815, 2816, 3005, 3017,
3075,
     3102, 3239, 3313, 3338, 3399, 3483, 3484, 3518, 3571, 3578,
3650,
     3806, 3819, 3886, 3894, 3925, 3941, 3978, 4037, 4145, 4206,
4298,
     4320, 4328, 4417, 4515, 4526, 4664, 4774, 4822, 4908, 5038,
5107,
     5226, 5231, 5249, 5369, 5446, 5604, 5655, 5781, 5825, 5838,
5854,
     5912, 5942, 6091, 6109, 6177, 6215, 6432, 6456, 6489, 6520,
6524,
     6557, 6562, 6706, 6721, 6803, 6968, 6974, 6992, 7096, 7233,
7276,
     7321, 7398, 7422, 7440, 7544, 7564, 7767, 7826, 7865, 7966,
7977,

```

```

7980, 8020, 8096, 8105, 8125, 8194, 8202, 8274, 8353, 8370,
8375,
8410, 8477, 8513, 8520, 8601, 8712, 8713, 8759, 8848, 8931,
8946,
8997, 9062, 9122, 9155, 9174, 9208, 9213, 9230, 9299, 9306,
9352,
9373, 9386, 9531, 9596, 9708, 9748, 9755, 9785, 9819, 9825,
9943,
9948, 9949, 9981, 9992, 9994, 9995, 9996])

def plot_images(dataset, ids, cols):
    dataset = copy.deepcopy(dataset)
    dataset.transform = A.Compose([t for t in dataset.transform if not
isinstance(t, (A.Normalize, ToTensorV2))])
    rows = len(ids) // cols
    figure, ax = plt.subplots(nrows=rows, ncols=cols, figsize=(18,
12))
    for i, idx in enumerate(ids):
        image, _ = dataset[idx]
        ax.ravel()[i].imshow(image)
        ax.ravel()[i].set_axis_off()
    plt.tight_layout()
    plt.show()

plot_images(test_dataset, wrong_preds_idx[:20], 5)

```



В большинстве картинок на которых ошиблась модель качество изображения такое, что даже человеку сложно разобрать какой именно номер (в особенности иероглиф) изображен на картинке